

Spring Batch Workshop

Arnaud Cogoluègnes
Consultant at Zenika, co-author “Spring Batch in Action”

Overview

- This workshop highlights Spring Batch features
- Problem/solution approach
 - A few slides to cover the feature
 - A project to start from, just follow the TODOs
- Prerequisites :
 - Basics about Java and Java EE
 - Spring: dependency injection, enterprise support
- <https://github.com/acogoluegnes/Spring-Batch-Workshop>

Settings

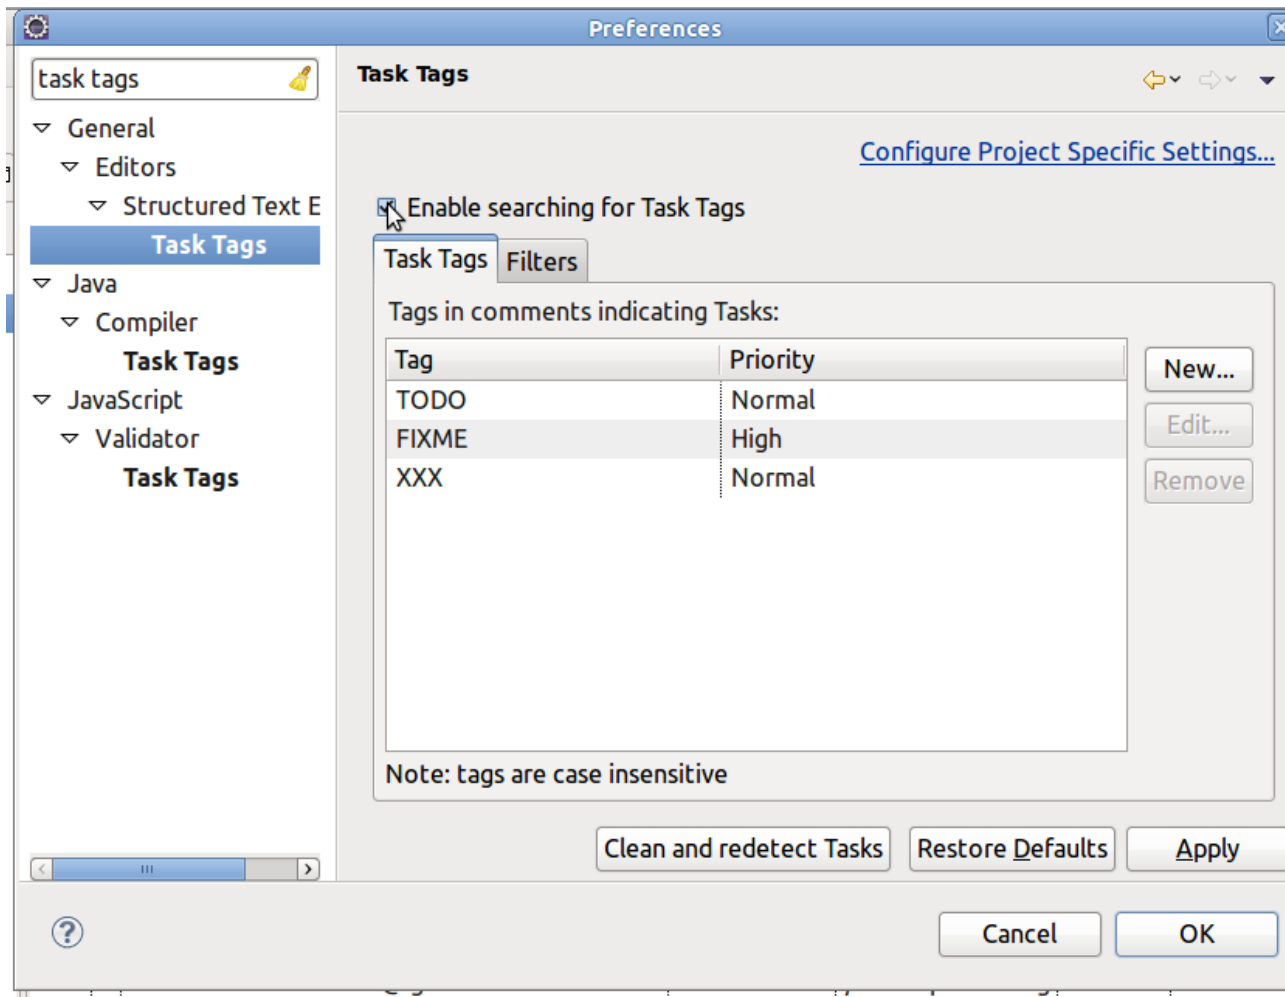
- Track the TODO in the *-start projects!
- It's easier with support from the IDE

The screenshot displays an IDE interface. On the left, a project explorer shows a tree structure with the following items: `hello-world-start`, `skip-solution`, `skip-start`, and `springbatch`. The main editor area is divided into tabs: `Problems`, `Javadoc`, `Declaration`, and `Tasks`. The `Tasks` tab is active, showing a list of 10 items. The list is organized into columns: `Description`, `Resource`, `Path`, `Location`, and `Type`. The items are numbered from 01 to 10, each with a specific description, resource, path, location, and type. The bottom status bar indicates the current file is `chunk-processing-start`.

	Description	Resource	Path	Location	Type
✓	TODO 01 fill in the list with String objects	StringItemRea	/chunk-processing-	line 22	Java Task
	TODO 02 remove items from the list until it's e	StringItemRea	/chunk-processing-	line 32	Java Task
	TODO 03 decorate the incoming item with star	StringItemProc	/chunk-processing-	line 19	Java Task
	TODO 04 output items one after the other with	StringItemWr	/chunk-processing-	line 25	Java Task
	TODO 05 declare Spring beans for the item read	chunk-process	/chunk-processing-	line 10	XML Task
	TODO 06 configure the job with the chunk-orie	chunk-process	/chunk-processing-	line 8	XML Task
	TODO 07 remove the @Ignore annotation on t	ChunkProcess	/chunk-processing-	line 21	Java Task
	TODO 08 run the job with the job launcher. Use	ChunkProcess	/chunk-processing-	line 32	Java Task
	TODO 09 check the returned JobExecution	ChunkProcess	/chunk-processing-	line 34	Java Task
	TODO 10 run the test! you should see the deco	ChunkProcess	/chunk-processing-	line 37	Java Task

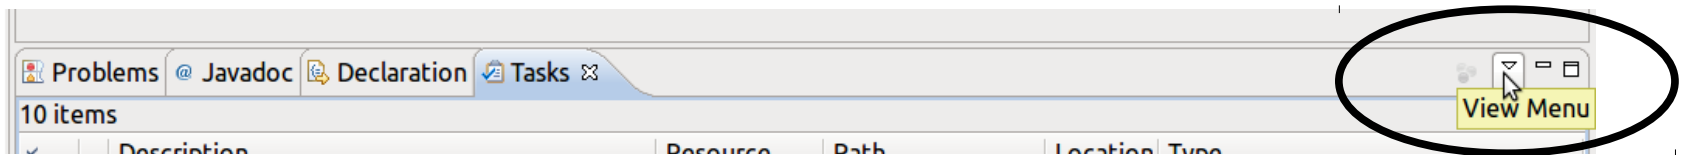
TODO with Eclipse

- Window > Preferences > “tasks tag” in filter

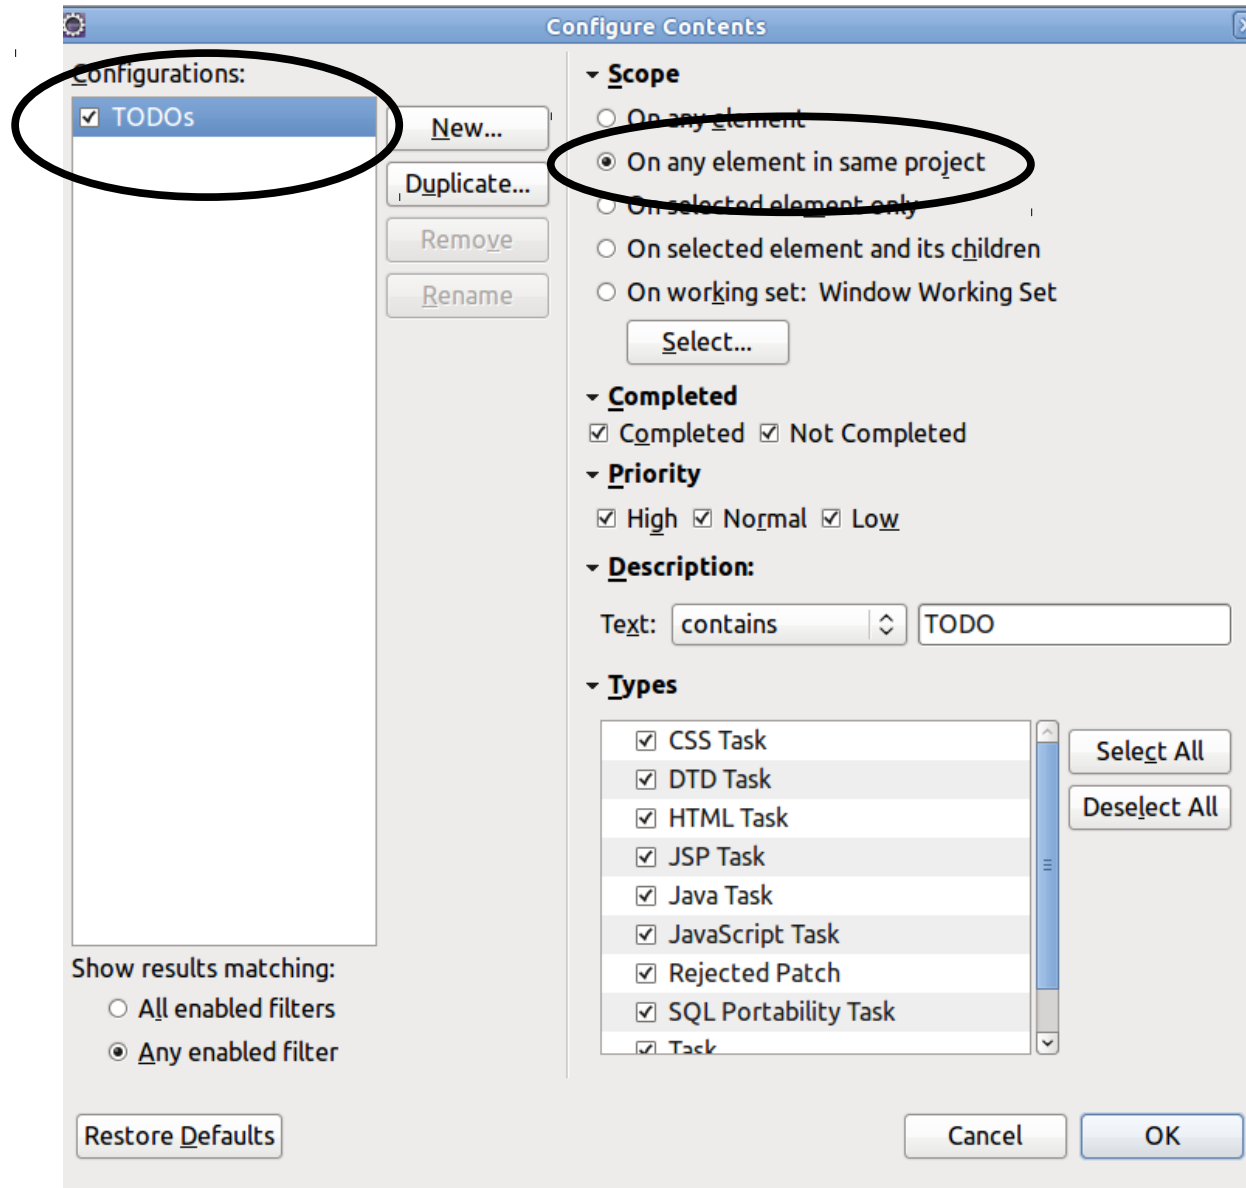


TODO with Eclipse

- Open the “Tasks” view and “configure contents”

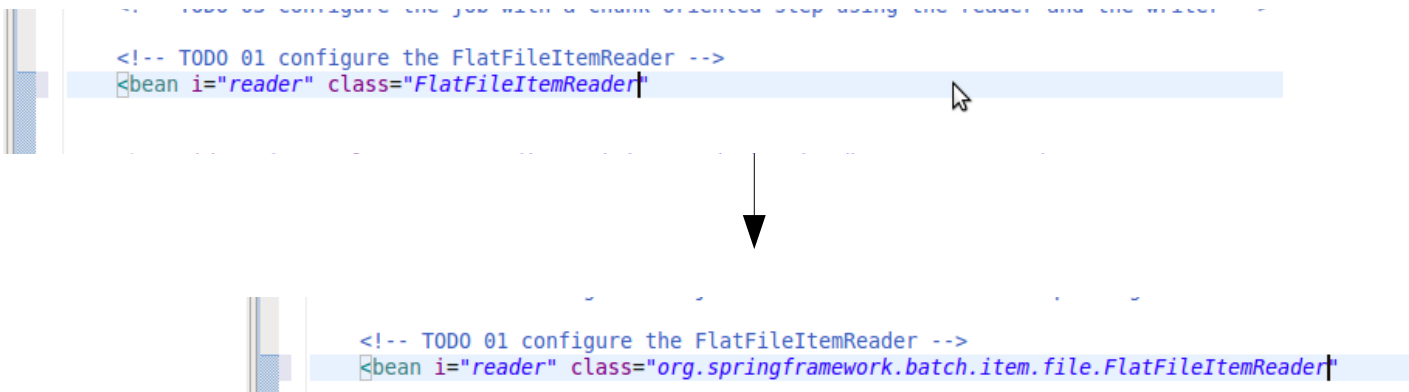


TODO with Eclipse



Spring support in IDE

- Spring support in IDE is a +
- e.g. code completion in SpringSource Tool Suite



Spring Batch overview

- Read – process – write large amounts of data, efficiently
- Ready-to-use components to read from/write to
 - Flat/XML files
 - Databases (JDBC, Hibernate, JPA, iBatis)
 - JMS queues
 - Emails
- Numerous extension points/hooks

Spring Batch overview

- Configuration to skip/retry items
- Execution metadata
 - Monitoring
 - Restart after failure
- Scaling strategies
 - Local/remote
 - Partitioning, remote processing

Hello World

- Problem: getting started with Spring Batch
- Solution: writing a simple “Hello world” job

Hello World

- A Spring Batch job is made of steps
- The Hello World job has one step
- The processing is implemented in a *Tasklet*

Hello World

- The Hello World Tasklet

```
public class HelloWorldTasklet implements Tasklet {  
  
    @Override  
    public RepeatStatus execute(  
        StepContribution contribution,  
        ChunkContext chunkContext) throws Exception {  
        System.out.println("Hello world!");  
        return RepeatStatus.FINISHED;  
    }  
}
```

Hello World

- The configuration of the Hello World job
 - Notice the `<batch />` namespace

```
<batch:job id="helloWorldJob">
  <batch:step id="helloWorldStep">
    <batch:tasklet>
      <bean class="com.zenika.workshop.springbatch.HelloWorldTasklet" />
    </batch:tasklet>
  </batch:step>
</batch:job>
```

Hello World

- Spring Batch needs some infrastructure beans
 - Let's use the typical test configuration

```
<bean id="transactionManager"
      class="o.s.b.support.transaction.ResourcelessTransactionManager" />

<bean id="jobRepository"
      class="o.s.b.core.repository.support.MapJobRepositoryFactoryBean" />

<bean id="jobLauncher"
      class="o.s.b.core.launch.support.SimpleJobLauncher">
  <property name="jobRepository" ref="jobRepository" />
</bean>
```

Hello World

- Let's test!

```
@RunWith(SpringJUnit4ClassRunner.class)
@ContextConfiguration("/hello-world-job.xml")
public class HelloWorldJobTest {

    @Autowired
    private Job job;

    @Autowired
    private JobLauncher jobLauncher;

    @Test public void helloWorld() throws Exception {
        JobExecution execution = jobLauncher.run(job, new JobParameters());
        assertEquals(ExitStatus.COMPLETED, execution.getExitStatus());
    }
}
```

Chunk processing

- Problem: processing large amounts of data efficiently
- Solution: using chunk processing

Chunk processing

- Batch jobs often read, process, and write items
- e.g.
 - Reading items from a file
 - Then processing (converting) items
 - Writing items to a database
- Spring Batch calls this “chunk processing”
 - a chunk = a set of items

Chunk processing

- Spring Batch
 - handles the iteration logic
 - uses a transaction for each chunk
 - lets you choose the chunk size
 - defines interfaces for each part of the processing

Chunk processing

- ItemReader
 - Reading ends when read() returns null

```
public interface ItemReader<T> {  
    T read() throws Exception, UnexpectedInputException,  
        ParseException, NonTransientResourceException;  
}
```

Chunk processing

- ItemProcessor
 - optional

```
public interface ItemProcessor<I, O> {  
    O process(I item) throws Exception;  
}
```

Chunk processing

- ItemWriting
 - Receive all the items of the chunk
 - Allows for batch update (more efficient)

```
public interface ItemWriter<T> {  
    void write(List<? extends T> items) throws Exception;  
}
```

Chunk processing

- Let's implement a (too?) simple chunk-oriented step!

Chunk processing

- The ItemReader

```
package com.zenika.workshop.springbatch;

import java.util.ArrayList;
import java.util.Arrays;
import java.util.List;
import org.springframework.batch.item.ItemReader;
import org.springframework.batch.item.NonTransientResourceException;
import org.springframework.batch.item.ParseException;
import org.springframework.batch.item.UnexpectedInputException;

public class StringItemReader implements ItemReader<String> {

    private List<String> list;

    public StringItemReader() {
        this.list = new ArrayList<String>(Arrays.asList("1","2","3","4","5","6","7"));
    }

    @Override
    public String read() throws Exception, UnexpectedInputException,
        ParseException, NonTransientResourceException {
        return !list.isEmpty() ? list.remove(0) : null;
    }
}
```

Chunk processing

- The ItemProcessor

```
package com.zenika.workshop.springbatch;

import org.springframework.batch.item.ItemProcessor;

public class StringItemProcessor implements ItemProcessor<String, String> {

    @Override
    public String process(String item) throws Exception {
        return "*** "+item+" ***";
    }

}
```


Chunk processing

- The ItemWriter

```
package com.zenika.workshop.springbatch;

import java.util.List;
import org.slf4j.Logger;
import org.slf4j.LoggerFactory;
import org.springframework.batch.item.ItemWriter;

public class StringItemWriter implements ItemWriter<String> {

    private static final Logger LOGGER =
        LoggerFactory.getLogger(StringItemWriter.class);

    @Override
    public void write(List<? extends String> items) throws Exception {
        for(String item : items) {
            LOGGER.info("writing "+item);
        }
    }
}
```

Chunk processing

- Configuring the job

```
<batch:job id="chunkProcessingJob">
  <batch:step id="chunkProcessingStep">
    <batch:tasklet>
      <batch:chunk reader="reader" processor="processor" writer="writer"
        commit-interval="3"
      />
    </batch:tasklet>
  </batch:step>
</batch:job>

<bean id="reader" class="com.zenika.workshop.springbatch.StringItemReader" />
<bean id="processor" class="com.zenika.workshop.springbatch.StringItemProcessor" />
<bean id="writer" class="com.zenika.workshop.springbatch.StringItemWriter" />
```

Chunk processing

- Do I always need to write my ItemReader/Processor/Writer?
- No, Spring Batch provides ready-to-use components for common datastores
 - Flat/XML files, databases, JMS, etc.
- You
 - Configure these components
 - Provides some logic
 - e.g. mapping a line with a domain object

Chunk processing

- Going further...
 - Reader/writer implementation for flat/XML files, database, JMS
 - Skipping items when something goes wrong
 - Listeners to react to the chunk processing

Flat file reading

- Problem: reading lines from a flat file and sending them to another source (e.g. database)
- Solution: using the FlatFileItemReader


Flat file reading

- Spring Batch has built-in support for flat files
 - Through the FlatFileItemReader for reading
- The FlatFileItemReader handles I/O
- 2 main steps:
 - Configuring the FlatFileItemReader
 - Providing a line – object mapping strategy

Flat file reading

- The usual suspects:

```
De-Anna,Raghunath,2010-03-04  
Susy,Hauerstock,2010-03-04  
Kiam,Whitehurst,2010-03-04  
Alecia, Van Holst,2010-03-04  
Hing,Senecal,2010-03-04
```



```
public class Contact {  
  
    private Long id;  
    private String firstname,lastname;  
    private Date birth;  
  
    (...)  
}
```

Flat file reading

- What do we need to read a flat file?
 - How to tokenize a line
 - How to map the line with a Java object
 - Where to find the file to read

Flat file reading

Tokenization

```
<bean id="reader"
      class="org.springframework.batch.item.file.FlatFileItemReader">
  <property name="lineMapper">
    <bean class="org.springframework.batch.item.file.mapping.DefaultLineMapper">
      <property name="lineTokenizer">
        <bean
          class="org.springframework.batch.item.file.transform.DelimitedLineTokenizer">
            <property name="names" value="firstname,lastname,birth" />
          </bean>
        </property>
        <property name="fieldSetMapper">
          <bean class="com.zenika.workshop.springbatch.ContactFieldSetMapper" />
        </property>
      </bean>
    </property>
    <property name="resource" value="classpath:contacts.csv" />
  </bean>
```

File to read

Line – object mapping

Flat file reading

- A FieldSetMapper to map a line with an object
- More about business logic, so typically implemented by developer
 - Spring Batch provides simple implementations

Flat file reading

```
package com.zenika.workshop.springbatch;

import org.springframework.batch.item.file.mapping.FieldSetMapper;
import org.springframework.batch.item.file.transform.FieldSet;
import org.springframework.validation.BindException;

public class ContactFieldSetMapper implements FieldSetMapper<Contact> {

    @Override
    public Contact mapFieldSet(FieldSet fieldSet) throws BindException {
        return new Contact(
            fieldSet.readString("firstname"),
            fieldSet.readString("lastname"),
            fieldSet.readDate("birth", "yyyy-MM-dd")
        );
    }
}
```

Flat file reading

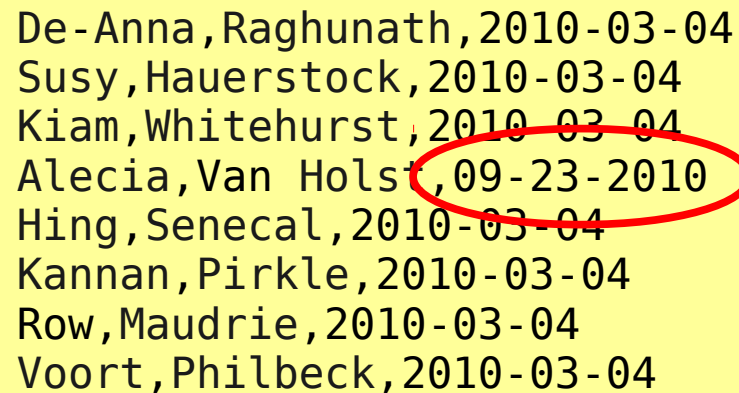
- Going further...
 - FlatFileItemWriter to write flat file
 - Fixed-length format (different tokenizer)
 - Skipping badly formatted lines

Skip

- Problem: my job fails miserably because of a tiny error in my input file
- Solution: skipping lines without failing the whole execution

Skip

- Skipping lines is sometimes acceptable



```
De-Anna,Raghunath,2010-03-04  
Susy,Hauerstock,2010-03-04  
Kiam,Whitehurst,2010-03-04  
Alecia, Van Holst,09-23-2010  
Hing,Senecal,2010-03-04  
Kannan,Pirkle,2010-03-04  
Row,Maudrie,2010-03-04  
Voort,Philbeck,2010-03-04
```

Skip

- Skip in Spring Batch
 - Choose the exceptions to skip
 - Set the max number of items to skip

```
<batch:job id="skipJob">
  <batch:step id="skipStep">
    <batch:tasklet>
      <batch:chunk reader="reader" writer="writer" commit-interval="3"
        skip-limit="10">
        <batch:skippable-exception-classes>
          <batch:include
            class="org.springframework.batch.item.file.FlatFileParseException"/>
        </batch:skippable-exception-classes>
      </batch:chunk>
    </batch:tasklet>
  </batch:step>
</batch:job>
```

Skip

- Going further...
 - Logging skipped items with a SkipListener
 - Setting a custom SkipPolicy

Dynamic job parameters

- Problem: passing values to the configuration when launching a job
- Solution: using job parameters and late binding

Dynamic job parameters

- Use case: providing a input file dynamically to the item reader

```
JobParameters jobParameters = new JobParametersBuilder()
    .addString("input.file", "file:./input/contacts-01.txt")
    .toJobParameters();
JobExecution execution = jobLauncher.run(job, jobParameters);
```

```
<bean id="reader"
      class="org.springframework.batch.item.file.FlatFileItemReader"
      scope="step">
  <property name="resource" value="#{jobParameters['input.file']}" />
  (...)
</bean>
```

Dynamic job parameters

- Going further...
 - Checking other available variables in an expression

JDBC paging

- Problem: reading large result sets from the database with a stable memory footprint
- Solution: using the `JdbcPagingItemReader`, which uses paging to handle large result sets

JDBC paging

```
<bean id="reader"
      class="org.springframework.batch.item.database.JdbcPagingItemReader">
  <property name="dataSource" ref="dataSource" />
  <property name="pageSize" value="10" />
  <property name="queryProvider">
    <bean class="o.s.b.item.database.support.SqlPagingQueryProviderFactoryBean">
      <property name="dataSource" ref="dataSource" />
      <property name="selectClause"
        value="select id,firstname,lastname,birth" />
      <property name="fromClause" value="from contact" />
      <property name="sortKey" value="id" />
    </bean>
  </property>
  <property name="rowMapper">
    <bean class="com.zenika.workshop.springbatch.ContactRowMapper" />
  </property>
</bean>
```

JDBC paging

- By paging, you send multiple queries to the database
- Alternative: cursor-based item reader
 - Spring Batch “streams” the result set from the DB
 - Only one query
- Paging always works, cursor-based reader depends on driver implementation

JDBC paging

- Going further...
 - Paging readers for Hibernate, JPA, iBatis
 - Cursor-based readers

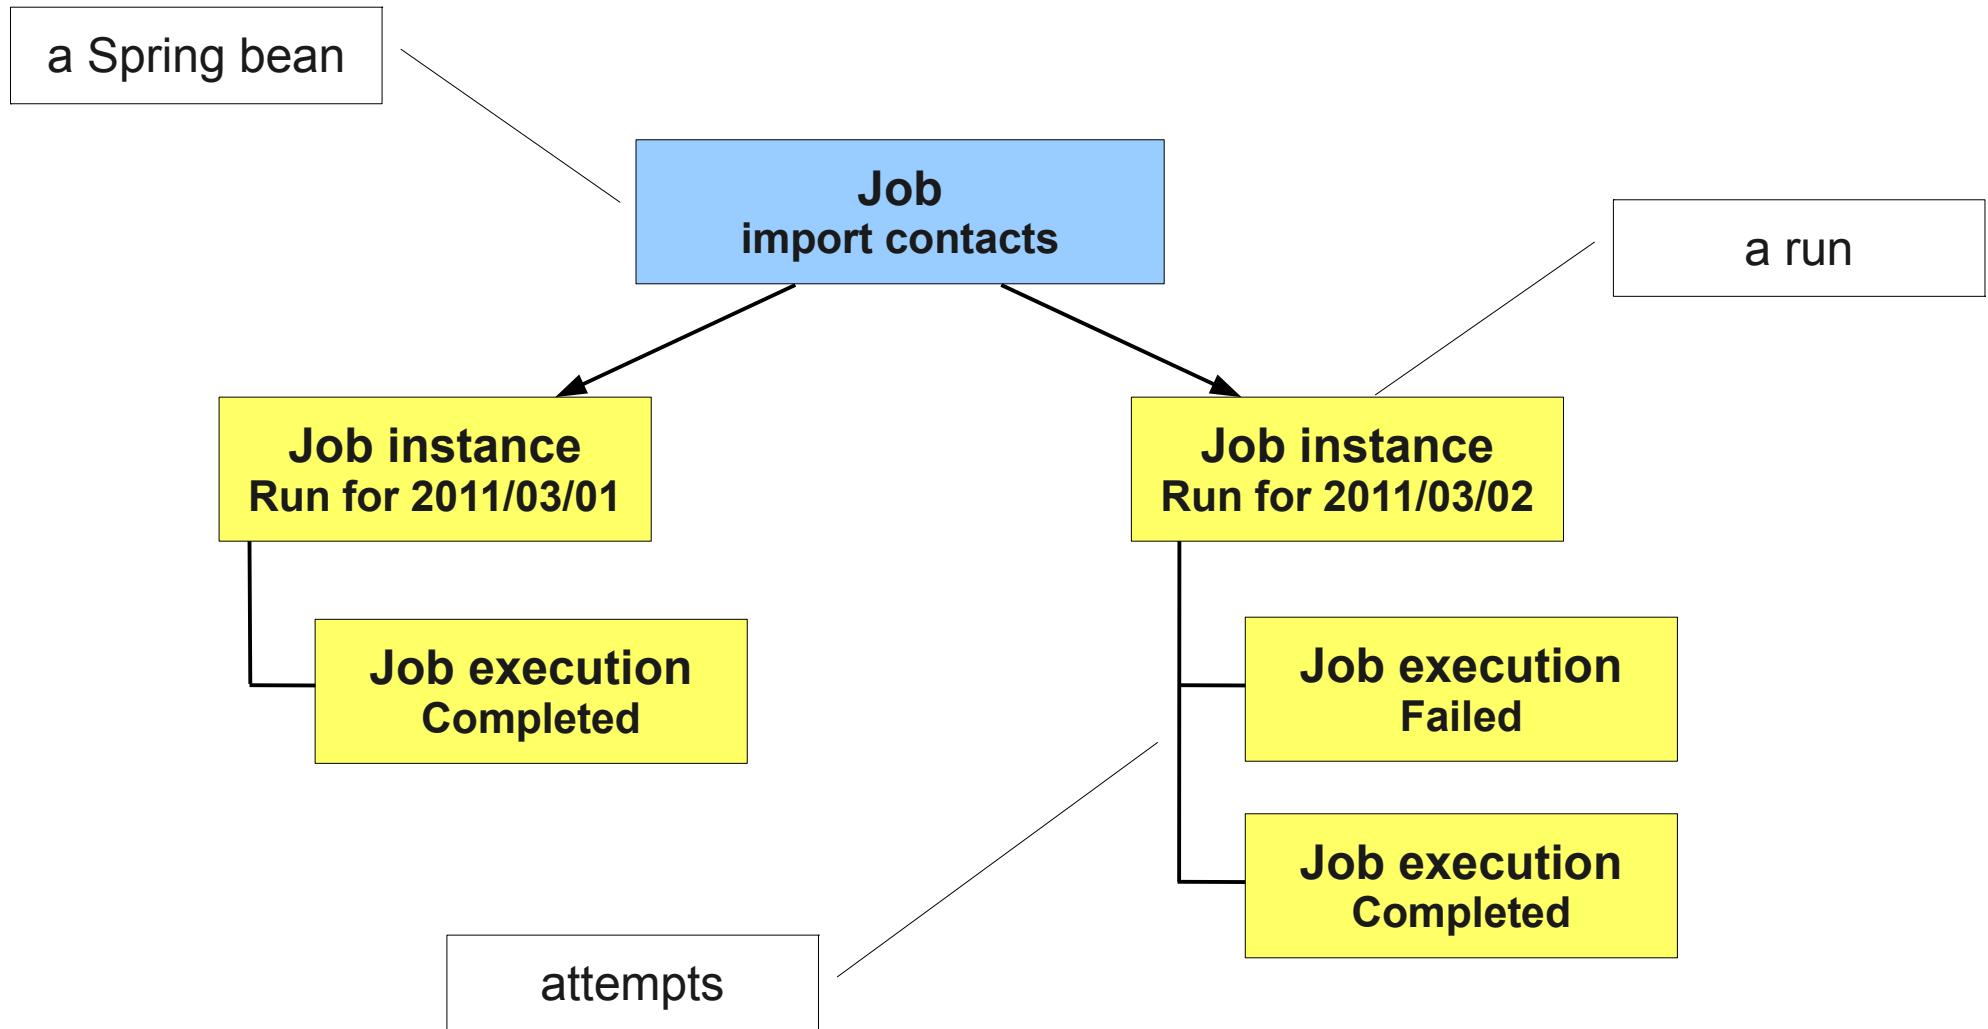
Execution metadata

- Problem: monitoring the execution of batch jobs
- Solution: letting Spring Batch storing execution metadata in a database

Execution metadata

- Spring Batch keeps track of batch execution
- Enables:
 - Monitoring by querying metadata tables
 - Restarting after a failure

Execution metadata



Execution metadata

- How to define a job instance?
- Thanks to job parameters
 - They define the identity of the job instance

Execution metadata

- Metadata are stored in a database
 - In-memory implementation for test/development
- Monitoring tools can query metadata tables
 - e.g. Spring Batch Admin

Execution metadata

- Going further...
 - Spring Batch Admin set-up
 - JobExplorer and JobOperator interfaces
 - Spring JMX support

Scheduling

- Problem: scheduling a job to execute periodically
- Solution: using the scheduling support in Spring

Scheduling

```
public class ImportLauncher {  
  
    public void launch() throws Exception {  
        JobExecution exec = jobLauncher.run(  
            job,  
            new JobParametersBuilder()  
                .addLong("time", System.currentTimeMillis())  
                .toJobParameters()  
        );  
    }  
}
```

```
<bean id="importLauncher"  
      class="com.zenika.workshop.springbatch.ImportLauncher" />  
  
<task:scheduled-tasks>  
    <task:scheduled ref="importLauncher" method="launch"  
                    fixed-delay="1000" />  
</task:scheduled-tasks>
```

A "cron" attribute is available

Scheduling

- Going further...
 - Threading settings in Spring Scheduler
 - Spring support for Quartz

Item processor

- Problem: I want to add some business logic before writing the items I just read
- Solution: use an item processor to process/convert read items before sending them to the item writer

Item processor

- Use case:
 - Reading contacts from a flat file
 - Registering them into the system
 - Writing the registration confirmations to the database



Business logic

Item processor

```
public interface ItemProcessor<I, O> {  
    O process(I item) throws Exception;  
}
```

Item processor

- Delegate to business service

```
public class ContactItemProcessor implements
    ItemProcessor<Contact, RegistrationConfirmation> {

    private RegistrationService registrationService;

    @Override
    public RegistrationConfirmation process(Contact item)
        throws Exception {
        return registrationService.process(item);
    }
}
```

Item processor

- Register the item processor on the step

```
<batch:job id="itemProcessorJob">
  <batch:step id="itemProcessorStep">
    <batch:tasklet>
      <batch:chunk reader="reader" processor="processor"
        writer="writer" commit-interval="3"/>
    </batch:tasklet>
  </batch:step>
</batch:job>

<bean id="registrationService"
  class="com.zenika.workshop.springbatch.RegistrationService" />

<bean id="processor"
  class="com.zenika.workshop.springbatch.ContactItemProcessor">
  <property name="registrationService" ref="registrationService" />
</bean>
```

Item processor

- Going further...
 - Available ItemProcessor implementations

Logging skipped items

- Problem: logging skipped items
- Solution: using a skip listener

Logging skipped items

- 2 steps:
 - Writing the skip listener (and the logging code)
 - Registering the listener on the step

Logging skipped items

- Writing the skip listener

```
package com.zenika.workshop.springbatch;

import org.slf4j.Logger;
import org.slf4j.LoggerFactory;
import org.springframework.batch.core.listener.SkipListenerSupport;

public class Slf4jSkipListener<T,S> extends SkipListenerSupport<T, S> {

    private static final Logger LOG = LoggerFactory.getLogger(
        Slf4jSkipListener.class);

    @Override
    public void onSkipInRead(Throwable t) {
        LOG.warn("skipped item: {}",t.toString());
    }

}
```

Logging skipped items

- Registering the skip listener

```
<batch:job id="loggingSkippedItemsJob">
  <batch:step id="loggingSkippedItemsStep">
    <batch:tasklet>
      <batch:chunk reader="reader" writer="writer" commit-interval="3"
        skip-limit="10">
        <batch:skippable-exception-classes>
          <batch:include
            class="org.springframework.batch.item.file.FlatFileParseException"/>
          </batch:skippable-exception-classes>
        </batch:chunk>
        <batch:listeners>
          <batch:listener ref="skipListener" />
        </batch:listeners>
      </batch:tasklet>
    </batch:step>
  </batch:job>

<bean id="skipListener" class="com.zenika.workshop.springbatch.Slf4jSkipListener" />
```

Logging skipped items

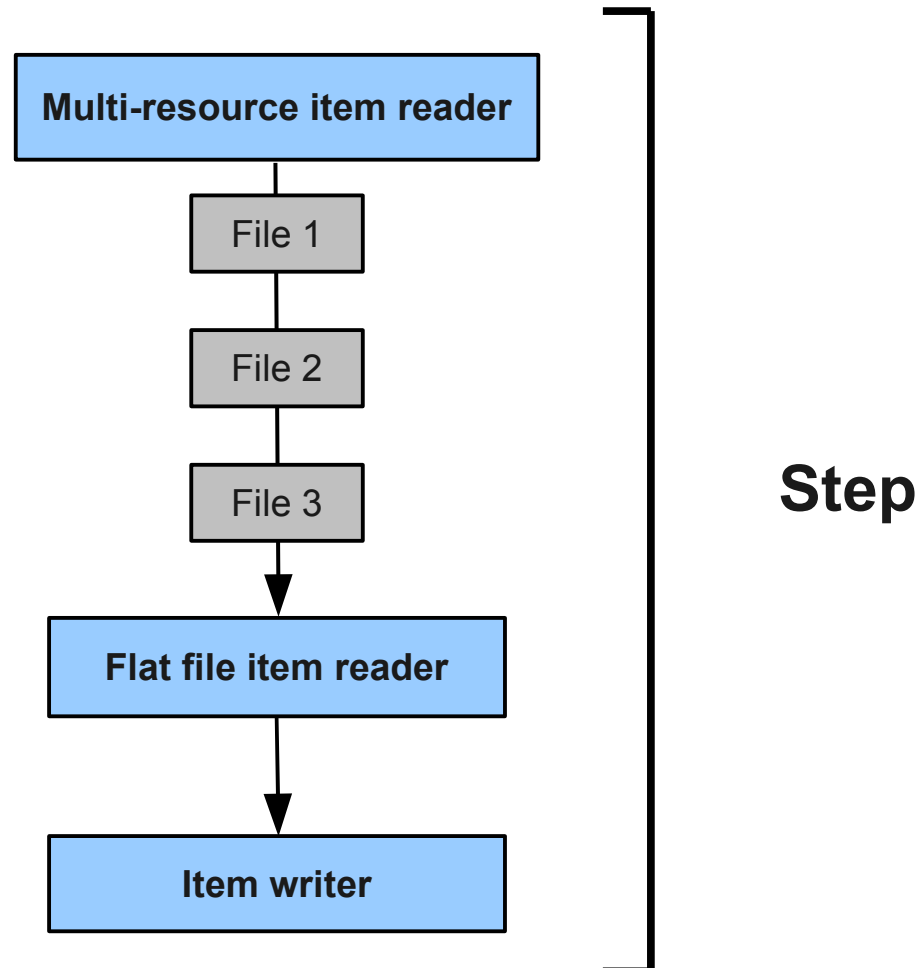
- Going further...
 - Other listeners in Spring Batch
 - ChunkListener, Item(Read/Process/Write)Listener, ItemStream, StepExecutionListener, JobExecutionListener

File reading partitioning

- Problem: I have multiple input files and I want to process them in parallel
- Solution: use partitioning to parallelize the processing on multiple threads

File reading partitioning

- No partitioning (one file after the other)

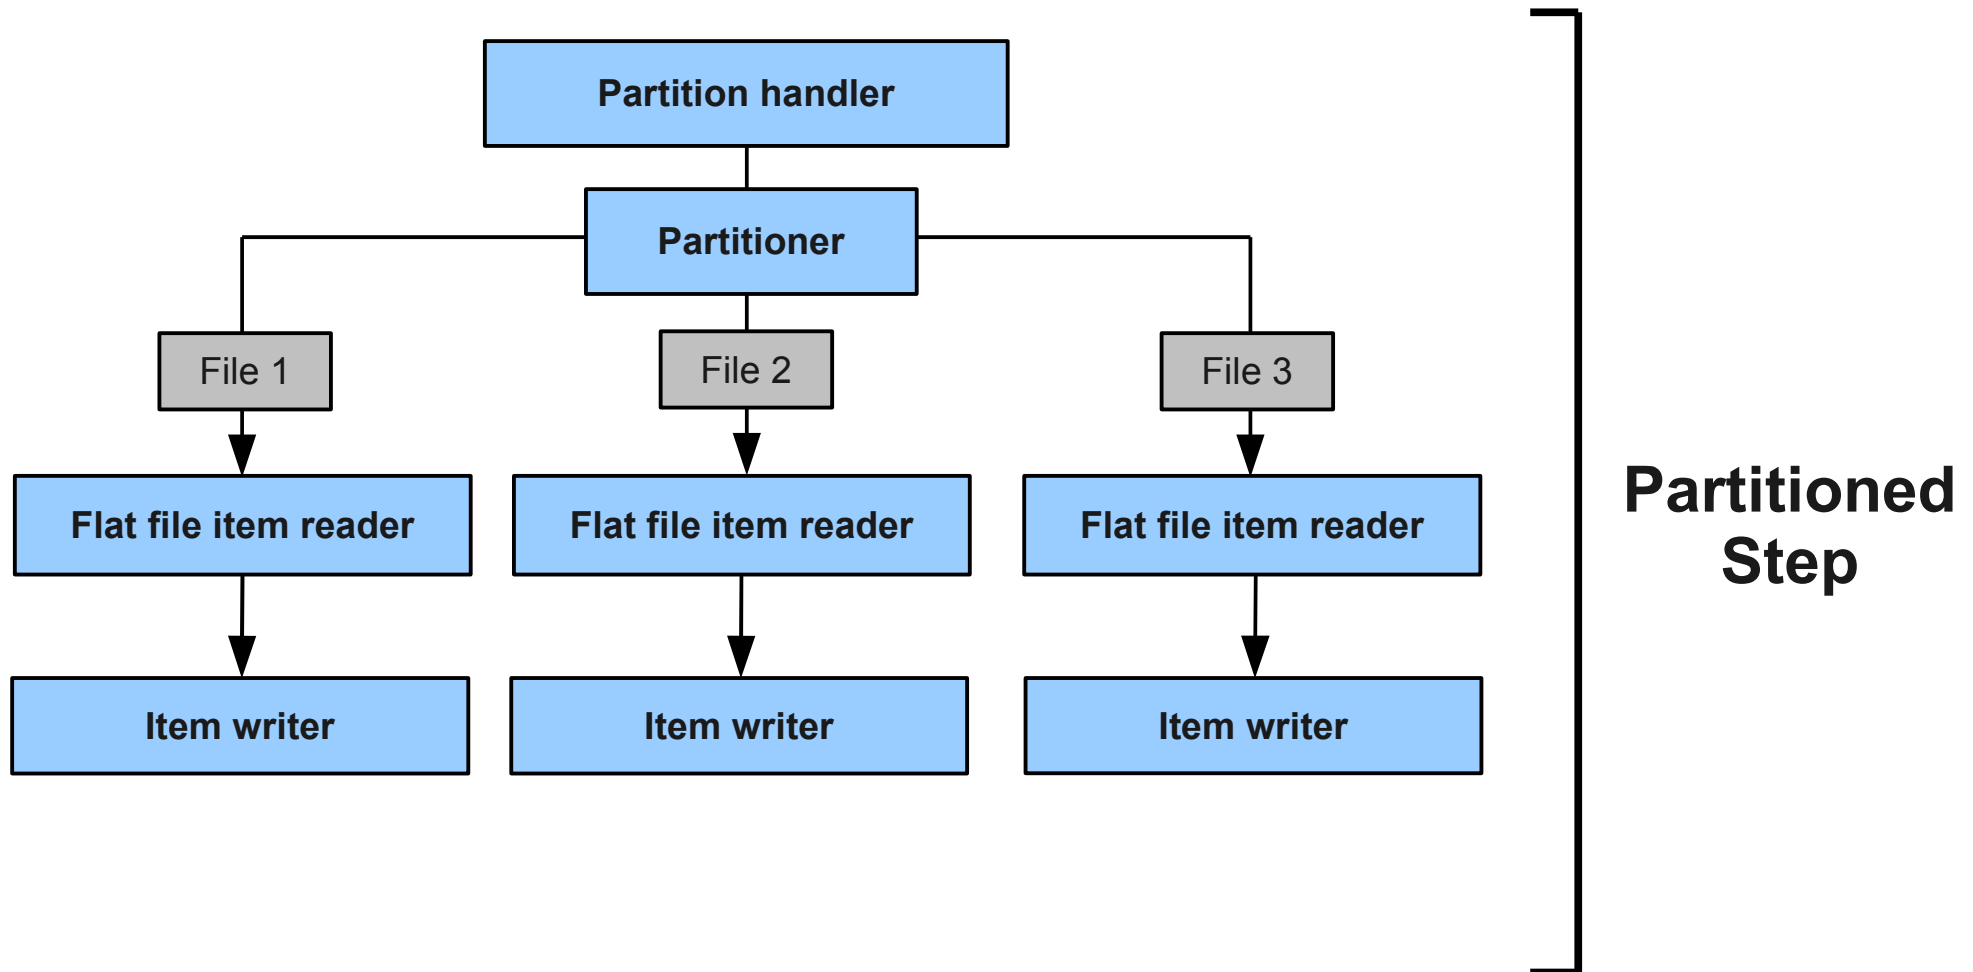


File reading partitioning

- Partitioning principle in Spring Batch:
 - Partition the data
 - e.g. one input file = one partition
 - Execute the partition in a dedicated step
- Partitioning is easy to set up but need some knowledge about the data
- Partition handler implementation
 - Multi-threaded
 - Spring Integration

File reading partitioning

- Multi-threaded partitioning



File reading partitioning

- Partitioner for input files

```
<bean id="partitioner"  
      class="o.s.b.core.partition.support.MultiResourcePartitioner">  
  <property name="resources"  
            value="file:./src/main/resources/input/*.txt" />  
</bean>
```

- Set a context for the steps to run

```
<bean id="reader"  
      class="org.springframework.batch.item.file.FlatFileItemReader"  
      scope="step">  
  (...)   
  <property name="resource" value="#{stepExecutionContext['fileName']}" />  
</bean>
```


File reading partitioning

- Using the multi-threaded partition handler

```
<batch:job id="fileReadingPartitioningJob">
  <batch:step id="partitionedStep" >
    <batch:partition step="readWriteContactsPartitionedStep"
                     partitioner="partitioner">
      <batch:handler task-executor="taskExecutor" />
    </batch:partition>
  </batch:step>
</batch:job>

<batch:step id="readWriteContactsPartitionedStep">
  <batch:tasklet>
    <batch:chunk reader="reader" writer="writer" commit-interval="10" />
  </batch:tasklet>
</batch:step>
```

File reading partitioning

- Going further...
 - Spring Integration partition handler implementation
 - Other scaling approaches (parallel steps, remote chunking, multi-threaded step)

File dropping launching

- Problem: downloading files from a FTP server and processing them with Spring Batch
- Solution: use Spring Integration to poll the FTP server and trigger Spring Batch accordingly

File dropping launching

FTP Server



1. Polls
2. Downloads

Spring Integration

4. Triggers

Spring Batch

3. Copies locally



File dropping launching

- The launching code

The local copy

```
public class FileContactJobLauncher {  
    public void launch(File file) throws Exception {  
        JobExecution exec = jobLauncher.run(  
            job,  
            new JobParametersBuilder()  
                .addString("input.file", "file:"+file.getAbsolutePath())  
                .toJobParameters()  
        );  
    }  
}
```

File dropping launching

- Listening to the FTP server

```
<int:channel id="fileIn" />

<int-ftp:inbound-channel-adapter local-directory="file:./input"
    channel="fileIn" session-factory="ftpClientFactory"
    remote-directory="/" auto-create-local-directory="true">
    <int:poller fixed-rate="1000" />
</int-ftp:inbound-channel-adapter>

<bean id="ftpClientFactory"
    class="com.zenika.workshop.springbatch.integration.DefaultFtpSessionFactory">
    <property name="host" value="localhost"/>
    <property name="port" value="2222"/>
    <property name="username" value="admin"/>
    <property name="password" value="admin"/>
</bean>
```

File dropping launching

- Calling the launcher on an inbound message

```
<int:channel id="fileIn" />

<int:service-activator input-channel="fileIn">
  <bean class="com.zenika.workshop.springbatch.integration.FileContactJobLauncher">
    <property name="job" ref="fileDroppingLaunchingJob" />
    <property name="jobLauncher" ref="jobLauncher" />
  </bean>
</int:service-activator>
```

File dropping message

- Going further...
 - Checking Spring Integration connectors
 - Local file system, FTPS, SFTP, HTTP, JMS, etc.
 - Checking operations on messages
 - Filtering, transforming, routing, etc.